



MongoDB, Mongoose and Cloud Storage

Frank Walsh, Diarmuid O'Connor

Agenda

- Cloud Databases
- MongoDB
- Mongoose
- Mongo in the cloud



Databases in Enterprise Apps

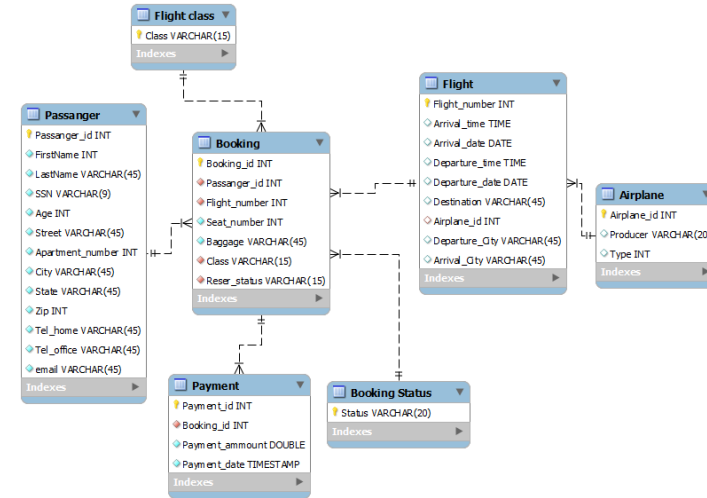
- Most data driven enterprise applications need a database
 - Persistence: storage of data
 - Concurrency: many applications sharing the data at once.
 - Integration: multiple systems using the same DB
- Enterprise Application DBs require backups, fail over, maintenance, capacity provisioning.
 - Traditionally handled by a Database Administrator (the DBA).



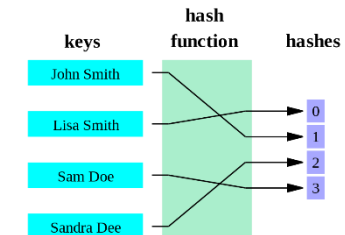
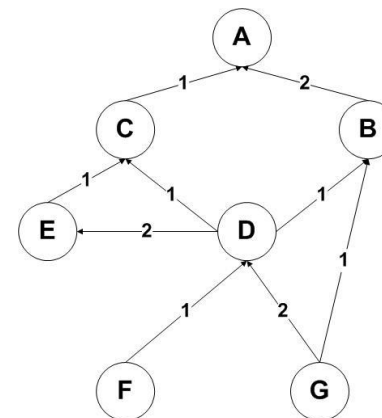
Structured & Unstructured Data

- Relational Databases:
 - Organise data into structured tables and rows
 - Relations have to be simple, they cannot contain any structure such as a nested record or a list
- In memory data structures
 - Much more varied structure
 - Lists, Queues, Stacks, Graphs, Hashing

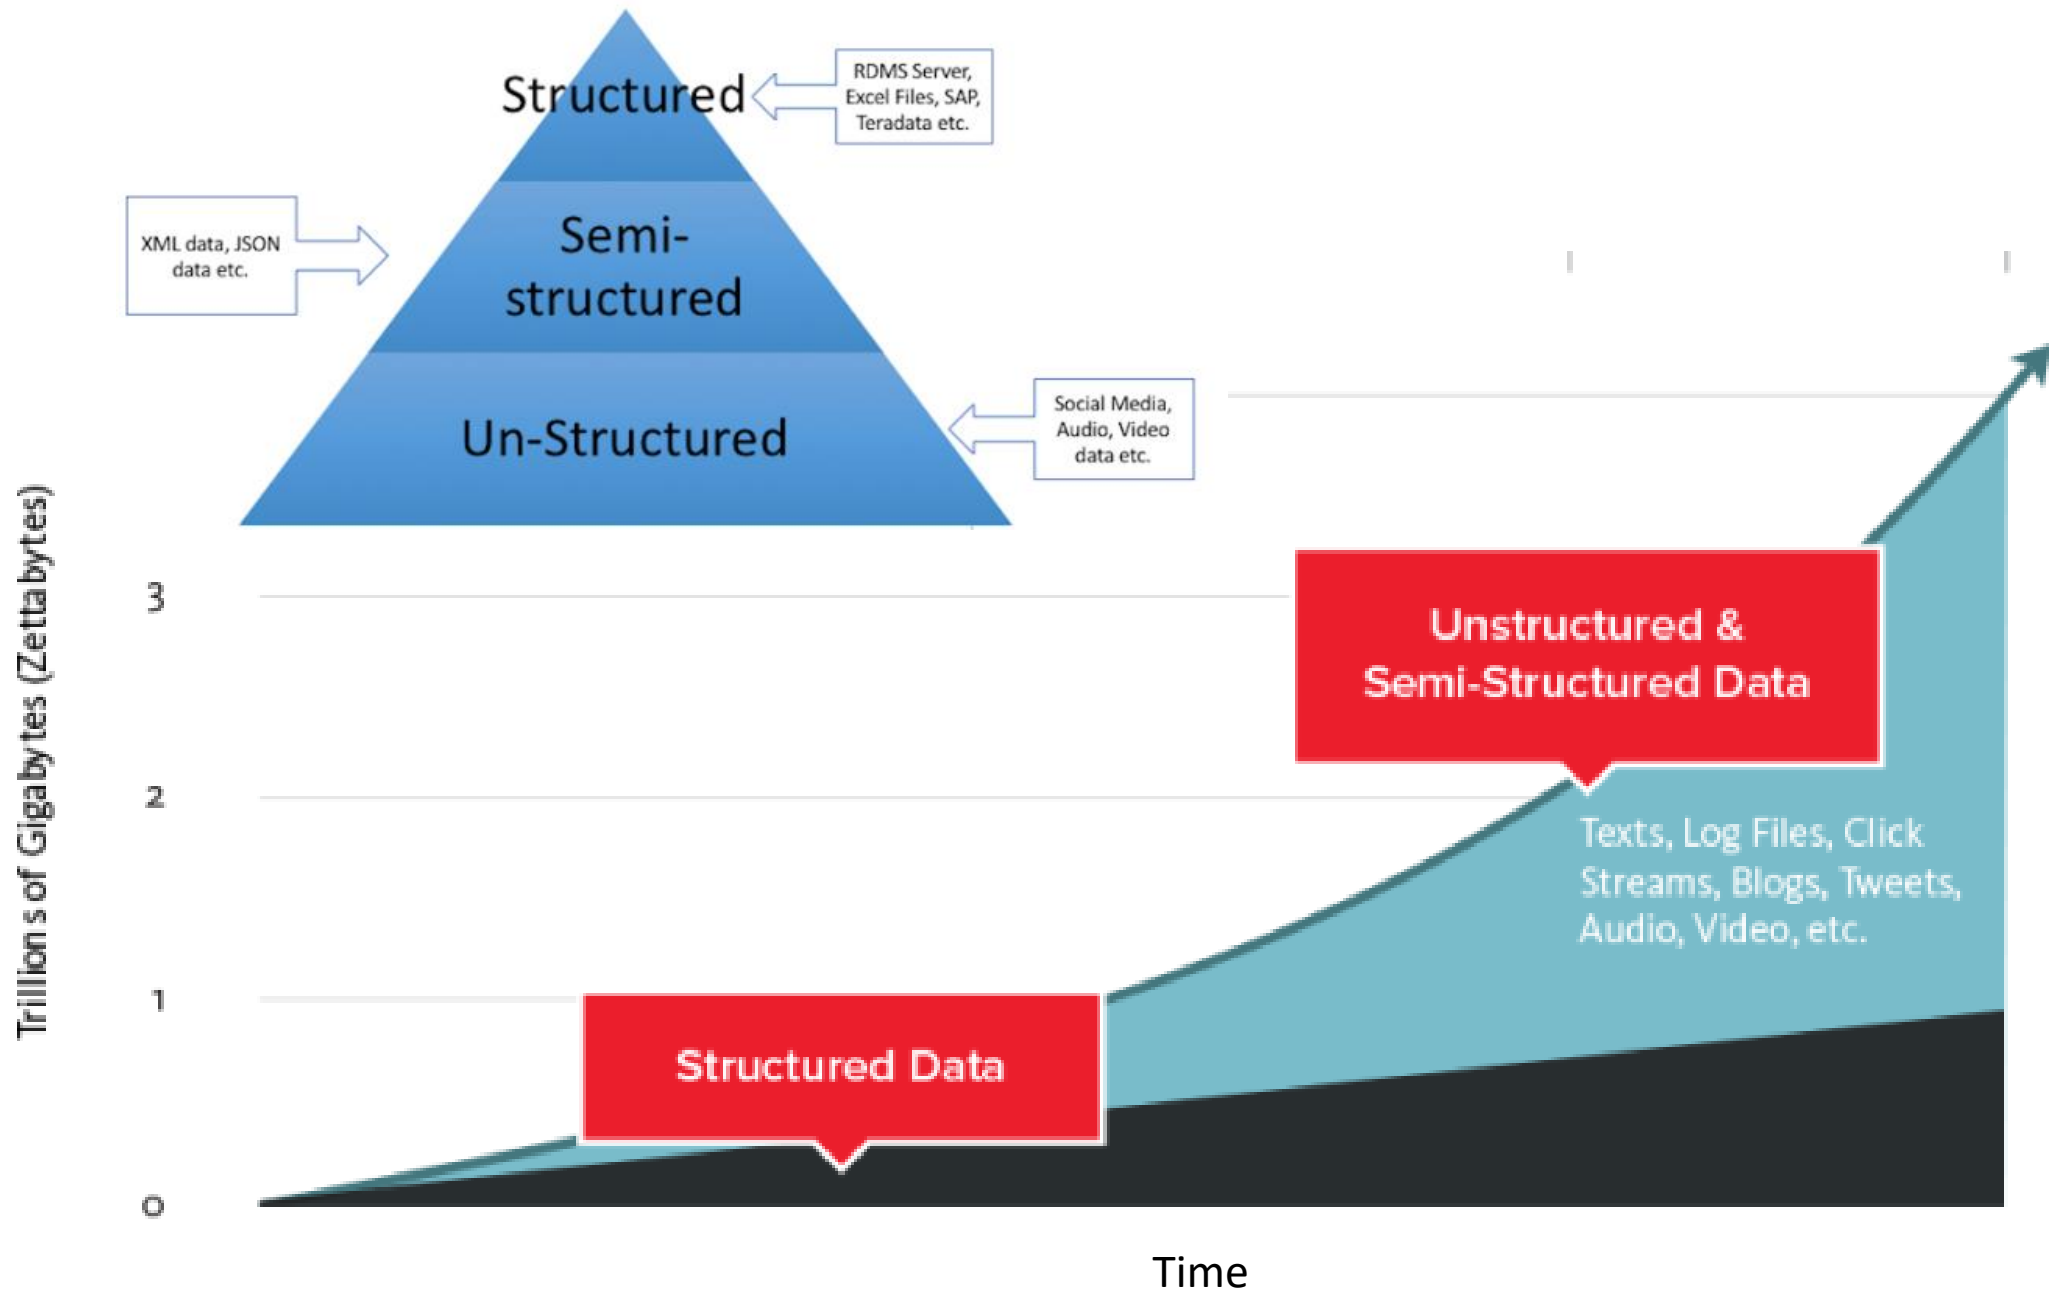
Relational Database



In Memory Data Structures



data format



Databases in the Cloud

- For some apps, a traditional relational database may not be the best fit
 - Organisations are capturing more data and processing it quicker – can be expensive/difficult on traditional DB
 - Traditionally, relational database is designed to run on a single machine in predictable environment
 - May be economic to run large data and computing loads on clusters.
 - Hard to estimate scaling requirements, particularly if it's a web app?
 - Are you going to do Data mining?
- One approach is to use the **Cloud** for your DB
 - Designed for scale
 - Can be outsourced so you don't have to deal with infrastructure requirements.



Cloud DB Advantages

- Removes Management costs
- Inherently scalable
- No need to define schemas(if NoSQL) etc.
- Lots of Cloud DB offerings out there
 - SQL based
 - NoSQL based
- If organisation policy/standards do not allow outsourcing:
 - Can host yourself, most NoSQL DBs are free.

Cloud Database Practices

- Drop Consistency
 - this makes distributed systems much easier to build
- Drop SQL and the relational model
 - simpler structures are easier to distribute:
 - key/value pairs
 - **structured documents**
 - **pseudo-tables**
 - tend to be schema-free, accepting data as-is
- Offer HTTP interfaces using XML or **JSON**
 - Web APIs!!!

Designing Distributed Data

- App data is not homogeneous
 - some kinds of data will be much larger
- consider using different databases for different requirements.
- user details,billing - needs consistency
 - require traditional database
- user data,content - needs partition tolerance
 - replicate to keep safe
- analytics,sessions - needs availability
 - "eventually consistent" is good enough